


GPT generation: predict, append, repeat

Training can score all next-token targets under a causal mask; inference extends the prefix one token at a time.



The generated token becomes part of the next input context.