

Scaled dot-product attention as matrix operations

QK^T creates pairwise relevance scores; masking removes illegal positions; softmax turns scores into weights; V supplies the returned information.

